



Statistical Rethinking

Week 9: Multilevel Models II Adventures in Covariance

Richard McElreath


Kinds of varying effects

- *Varying intercepts*: means differ by cluster
- *Varying slopes*: effects of predictors vary by cluster
- Any parameter can be made into a varying effect
 - (1) split into vector of parameters by cluster
 - (2) define population distribution



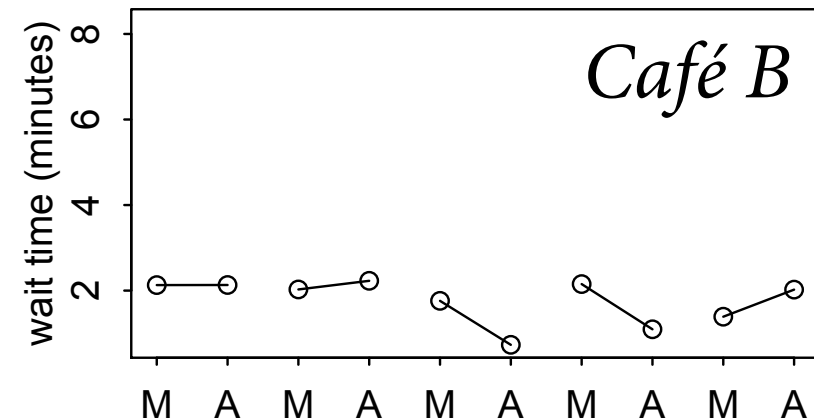

Varying slopes

- Why varying slopes?
 - drugs affect people differently
 - after school programs don't work for everyone
 - not every unit has same relationship to predictor
 - variation is important, whether for intervention or inference
- *Average* effect misleading?
- Pooling, shrinkage, mnesia

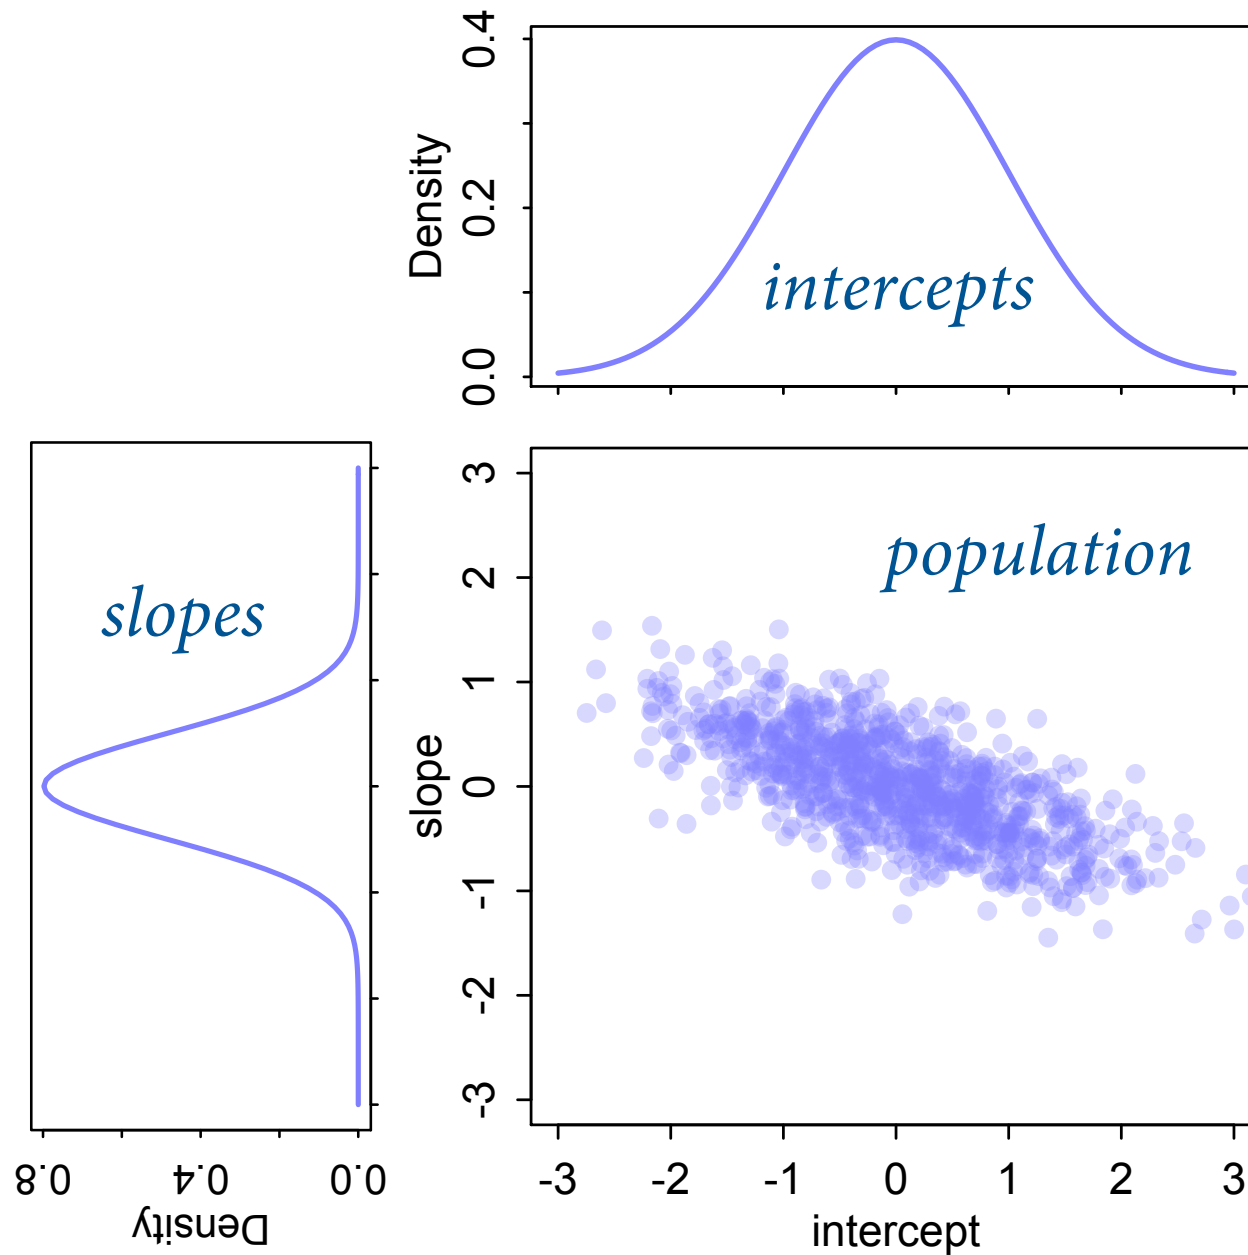


Café Robot

- Robot programmed to visit cafés, order coffee, record wait time
- Visits in *morning* and *afternoon*
- Intercepts: avg morning wait
- Slopes: avg difference btw afternoon and morning
- Are intercepts and slopes related?
 - Yes => pooling across parameter types!


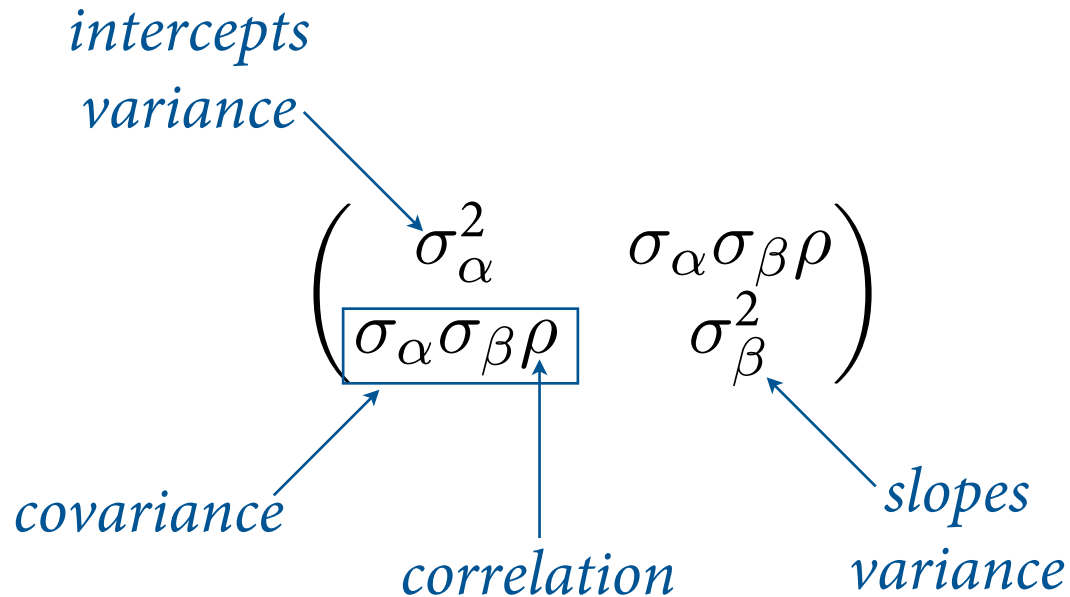


Population of Cafés




Population of Cafés

- 2-dimensional Gaussian distribution
 - vector of means
 - variance-covariance matrix



Simulated Cafés



20 cafés

5 days

morning & afternoon

200 observations

Varying slopes model

$$W_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha_{\text{CAFÉ}[i]} + \beta_{\text{CAFÉ}[i]} M_i$$

$$\begin{bmatrix} \alpha_{\text{CAFÉ}} \\ \beta_{\text{CAFÉ}} \end{bmatrix} \sim \text{MVNormal} \left(\begin{bmatrix} \alpha \\ \beta \end{bmatrix}, \mathbf{S} \right)$$

$$\mathbf{S} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \mathbf{R} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix}$$

$$\alpha \sim \text{Normal}(0, 10)$$

$$\beta \sim \text{Normal}(0, 10)$$

$$\sigma \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\alpha \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\beta \sim \text{HalfCauchy}(0, 1)$$

$$\mathbf{R} \sim \text{LKJcorr}(2)$$

varying intercepts —————

varying slopes —————

$$W_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha_{\text{CAFÉ}[i]} + \beta_{\text{CAFÉ}[i]} M_i$$

$$\begin{bmatrix} \alpha_{\text{CAFÉ}} \\ \beta_{\text{CAFÉ}} \end{bmatrix} \sim \text{MVNormal} \left(\begin{bmatrix} \alpha \\ \beta \end{bmatrix}, \mathbf{S} \right)$$

$$\mathbf{S} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \mathbf{R} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix}$$

$$\alpha \sim \text{Normal}(0, 10)$$

$$\beta \sim \text{Normal}(0, 10)$$

$$\sigma \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\alpha \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\beta \sim \text{HalfCauchy}(0, 1)$$

$$\mathbf{R} \sim \text{LKJcorr}(2)$$

$$W_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha_{\text{CAFÉ}[i]} + \beta_{\text{CAFÉ}[i]} M_i$$

multivariate prior $\longrightarrow \begin{bmatrix} \alpha_{\text{CAFÉ}} \\ \beta_{\text{CAFÉ}} \end{bmatrix} \sim \text{MVNormal} \left(\begin{bmatrix} \alpha \\ \beta \end{bmatrix}, \mathbf{S} \right)$

$$\mathbf{S} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \mathbf{R} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix}$$

$$\alpha \sim \text{Normal}(0, 10)$$

$$\beta \sim \text{Normal}(0, 10)$$

$$\sigma \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\alpha \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\beta \sim \text{HalfCauchy}(0, 1)$$

$$\mathbf{R} \sim \text{LKJcorr}(2)$$

$$W_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha_{\text{CAFÉ}[i]} + \beta_{\text{CAFÉ}[i]} M_i$$

pop avg intercept

*covariance
matrix*

pop avg slope

$$\begin{bmatrix} \alpha_{\text{CAFÉ}} \\ \beta_{\text{CAFÉ}} \end{bmatrix} \sim \text{MVNormal} \left(\begin{bmatrix} \alpha \\ \beta \end{bmatrix}, \mathbf{S} \right)$$

$$\mathbf{S} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \mathbf{R} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix}$$

$$\alpha \sim \text{Normal}(0, 10)$$

$$\beta \sim \text{Normal}(0, 10)$$

$$\sigma \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\alpha \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\beta \sim \text{HalfCauchy}(0, 1)$$

$$\mathbf{R} \sim \text{LKJcorr}(2)$$

Covariance matrix shuffle


- m -by- m covariance matrix requires estimating
 - m standard deviations (or variances)
 - $(m^2 - m)/2$ correlations (or covariances)
 - total of $m(m + 1)/2$ parameters
- Several ways specify priors
 - Conjugate: inverse-Wishart (`inv_wishart`)
 - inverse-Wishart cannot pull apart stddev and correlations
 - Better to decompose:

$$\mathbf{S} = \begin{pmatrix} \sigma_\alpha^2 & \rho\sigma_\alpha\sigma_\beta \\ \rho\sigma_\alpha\sigma_\beta & \sigma_\beta^2 \end{pmatrix} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \underbrace{\begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix}}_{\mathbf{R}} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix}$$

Matrixes are nice

- Matrix algebra just shortcuts for working with lists of numbers
- A few simple rules
- Can you make an omelet?
You can multiply matrixes.

Case	Singular	Plural
nominative	mātrīx	mātrīcēs
genitive	mātrīcis	mātrīcum
dative	mātrīcī	mātrīcibus
accusative	mātrīcem	mātrīcēs
ablative	mātrīce	mātrīcibus
vocative	mātrīx	mātrīcēs



$$\begin{pmatrix} a & b \\ c & d \end{pmatrix}$$

$$\begin{pmatrix} A & B \\ C & D \end{pmatrix}$$

$$\begin{pmatrix} a & b \\ c & d \end{pmatrix}$$

$$\begin{pmatrix} A & B \\ C & D \end{pmatrix}$$

$$\begin{pmatrix} a & b \\ c & d \end{pmatrix}$$

$$\begin{pmatrix} A \\ C \end{pmatrix}$$

B

D

Aa + Bc	

$$\begin{pmatrix} a & b \\ c & d \end{pmatrix}$$

$$\begin{pmatrix} A & B \\ C & D \end{pmatrix}$$

Aa + Bc	Ab + Bd
Ca + Dc	Cb + Dd

Matrixes are nice

$$\mathbf{S} = \begin{pmatrix} \sigma_{\alpha}^2 & \rho\sigma_{\alpha}\sigma_{\beta} \\ \rho\sigma_{\alpha}\sigma_{\beta} & \sigma_{\beta}^2 \end{pmatrix} = \begin{pmatrix} \sigma_{\alpha} & 0 \\ 0 & \sigma_{\beta} \end{pmatrix} \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix} \begin{pmatrix} \sigma_{\alpha} & 0 \\ 0 & \sigma_{\beta} \end{pmatrix} = \mathbf{SRS}$$

Matrixes are nice

$$\mathbf{S} = \begin{pmatrix} \sigma_\alpha^2 & \rho\sigma_\alpha\sigma_\beta \\ \rho\sigma_\alpha\sigma_\beta & \sigma_\beta^2 \end{pmatrix} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} = \mathbf{SRS}$$

$$\begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix}$$

$$\begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix}$$

?

Matrixes are nice

$$\mathbf{S} = \begin{pmatrix} \sigma_\alpha^2 & \rho\sigma_\alpha\sigma_\beta \\ \rho\sigma_\alpha\sigma_\beta & \sigma_\beta^2 \end{pmatrix} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} = \mathbf{SRS}$$

$$\begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix}$$

$$\begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \begin{pmatrix} \sigma_\alpha & \rho\sigma_\alpha \\ \rho\sigma_\beta & \sigma_\beta \end{pmatrix}$$

Matrixes are nice

$$\mathbf{S} = \begin{pmatrix} \sigma_\alpha^2 & \rho\sigma_\alpha\sigma_\beta \\ \rho\sigma_\alpha\sigma_\beta & \sigma_\beta^2 \end{pmatrix} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} = \mathbf{SRS}$$

$$\begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix}$$

$$\begin{pmatrix} \sigma_\alpha & \rho\sigma_\alpha \\ \rho\sigma_\beta & \sigma_\beta \end{pmatrix}$$

?

Matrixes are nice

$$\mathbf{S} = \begin{pmatrix} \sigma_\alpha^2 & \rho\sigma_\alpha\sigma_\beta \\ \rho\sigma_\alpha\sigma_\beta & \sigma_\beta^2 \end{pmatrix} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \begin{pmatrix} 1 & \rho \\ \rho & 1 \end{pmatrix} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} = \mathbf{SRS}$$

$$\begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix}$$

$$\begin{pmatrix} \sigma_\alpha & \rho\sigma_\alpha \\ \rho\sigma_\beta & \sigma_\beta \end{pmatrix} \begin{pmatrix} \sigma_\alpha^2 & \rho\sigma_\alpha\sigma_\beta \\ \rho\sigma_\alpha\sigma_\beta & \sigma_\beta^2 \end{pmatrix}$$

$$W_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha_{\text{CAFÉ}[i]} + \beta_{\text{CAFÉ}[i]} M_i$$

$$\begin{bmatrix} \alpha_{\text{CAFÉ}} \\ \beta_{\text{CAFÉ}} \end{bmatrix} \sim \text{MVNormal} \left(\begin{bmatrix} \alpha \\ \beta \end{bmatrix}, \mathbf{S} \right)$$

build cov matrix $\longrightarrow \mathbf{S} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \mathbf{R} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix}$

$$\alpha \sim \text{Normal}(0, 10)$$

$$\beta \sim \text{Normal}(0, 10)$$

$$\sigma \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\alpha \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\beta \sim \text{HalfCauchy}(0, 1)$$

$$\mathbf{R} \sim \text{LKJcorr}(2)$$

$$W_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha_{\text{CAFÉ}[i]} + \beta_{\text{CAFÉ}[i]} M_i$$

$$\begin{bmatrix} \alpha_{\text{CAFÉ}} \\ \beta_{\text{CAFÉ}} \end{bmatrix} \sim \text{MVNormal} \left(\begin{bmatrix} \alpha \\ \beta \end{bmatrix}, \mathbf{S} \right)$$

$$\mathbf{S} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \mathbf{R} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix}$$

$$\alpha \sim \text{Normal}(0, 10)$$

$$\beta \sim \text{Normal}(0, 10)$$

$$\sigma \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\alpha \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\beta \sim \text{HalfCauchy}(0, 1)$$

$$\mathbf{R} \sim \text{LKJcorr}(2)$$

fixed (non-adaptive) priors



$$W_i \sim \text{Normal}(\mu_i, \sigma)$$

$$\mu_i = \alpha_{\text{CAFÉ}[i]} + \beta_{\text{CAFÉ}[i]} M_i$$

$$\begin{bmatrix} \alpha_{\text{CAFÉ}} \\ \beta_{\text{CAFÉ}} \end{bmatrix} \sim \text{MVNormal} \left(\begin{bmatrix} \alpha \\ \beta \end{bmatrix}, \mathbf{S} \right)$$

$$\mathbf{S} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \mathbf{R} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix}$$

$$\alpha \sim \text{Normal}(0, 10)$$

$$\beta \sim \text{Normal}(0, 10)$$

$$\sigma \sim \text{HalfCauchy}(0, 1)$$

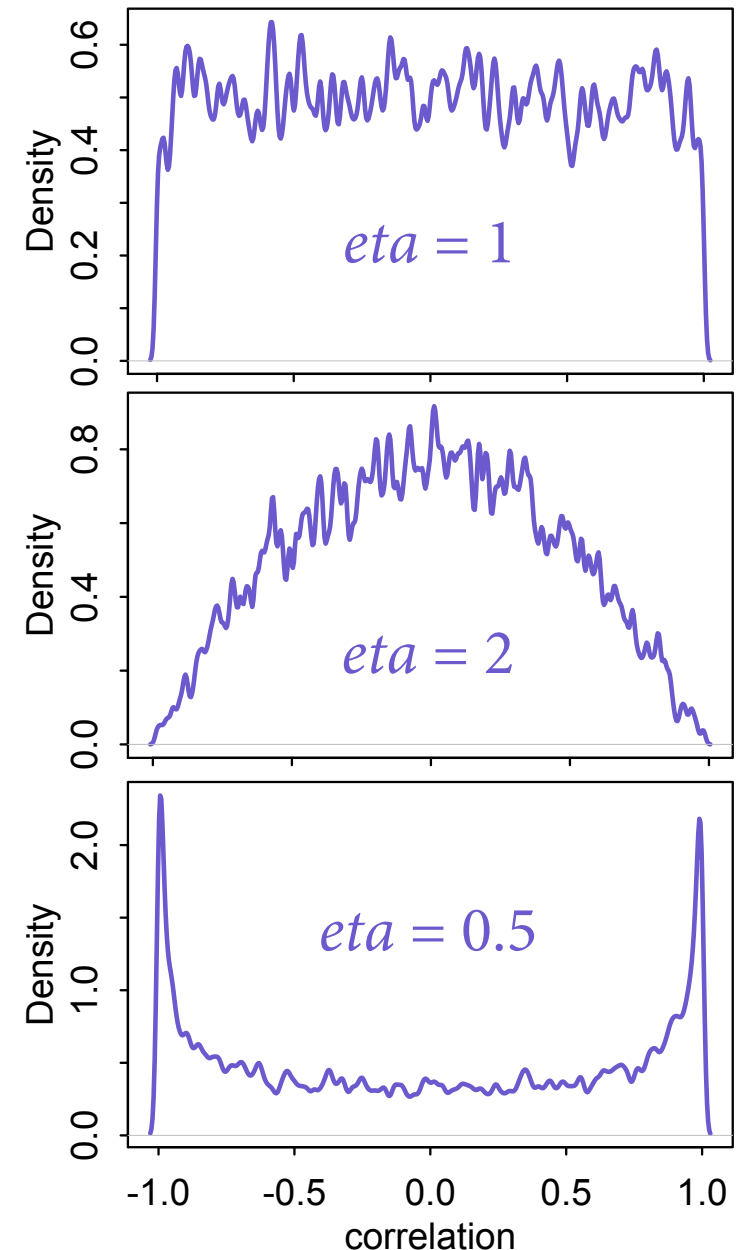
$$\sigma_\alpha \sim \text{HalfCauchy}(0, 1)$$

$$\sigma_\beta \sim \text{HalfCauchy}(0, 1)$$

correlation matrix prior → $\mathbf{R} \sim \text{LKJcorr}(2)$

LKJ Correlation prior

- After Lewandowski, Kurowicka, and Joe (LKJ) 2009
- One parameter, η , specifies concentration or dispersion from *identity matrix* (zero correlations)
 - $\eta = 1$, uniform correlation matrices
 - $\eta > 1$, stomps on extreme correlations
 - $\eta < 1$, elevates extreme correlations



Varying slopes estimation

```
m13.1 <- map2stan(  
  alist(  
    wait ~ dnorm( mu , sigma ),  
    mu <- a_cafe[cafe] + b_cafe[cafe]*afternoon,  
    c(a_cafe,b_cafe)[cafe] ~ dmvmnorm2(c(a,b),sigma_cafe,Rho),  
    a ~ dnorm(0,10),  
    b ~ dnorm(0,10),  
    sigma_cafe ~ dcauchy(0,2),  
    sigma ~ dcauchy(0,2),  
    Rho ~ dlkjcorr(2)  
  ) ,  
  data=d ,  
  iter=5000 , warmup=2000 , chains=2 )
```

Varying slopes estimation

```
m13.1 <- map2stan(  
  alist(  
    wait ~ dnorm( mu , sigma ),  
    mu <- a_cafe[cafe] + b_cafe[cafe]*afternoon,  
    c(a_cafe,b_cafe)[cafe] ~ dmvmnorm2(c(a,b),sigma_cafe,Rho),  
    a ~ dnorm(0,10),  
    b ~ dnorm(0,10),  
    sigma_cafe ~ dcauchy(0,2),  
    sigma ~ dcauchy(0,2),  
    Rho ~ dlkjcorr(2)  
  ) ,  
  data=d ,  
  iter=5000 , warmup=2000 , chains=2 )
```

Varying slopes estimation


```
m13.1 <- map2stan(  
  alist(  
    wait ~ dnorm( mu , sigma ),  
    mu <- a_cafe[cafe] + b_cafe[cafe]*afternoon,  
    c(a_cafe,b_cafe)[cafe] ~ dmnorm2(c(a,b),sigma_cafe,Rho),  
    a ~ dnorm(0,10),  
    b ~ dnorm(0,10),  
    sigma_cafe ~ dcauchy(0,2),  
    sigma ~ dcauchy(0,2),  
    Rho ~ dlkjcorr(2)  
  ) ,  
  data=d ,  
  iter=5000 , warmup=2000 , chains=2 )
```

Varying slopes estimation



```
m13.1 <- map2stan(  
  alist(  
    wait ~ dnorm( mu , sigma ),  
    mu <- a_cafe[cafe] + b_cafe[cafe]*afternoon,  
    c(a_cafe,b_cafe)[cafe] ~ dmvmnorm2(c(a,b),sigma_cafe,Rho),  
    a ~ dnorm(0,10),  
    b ~ dnorm(0,10),  
    sigma_cafe ~ dcauchy(0,2),  
    sigma ~ dcauchy(0,2),  
    Rho ~ dlkjcrr(2)  
  ) ,  
  data=d ,  
  iter=5000 , warmup=2000 , chains=2 )
```

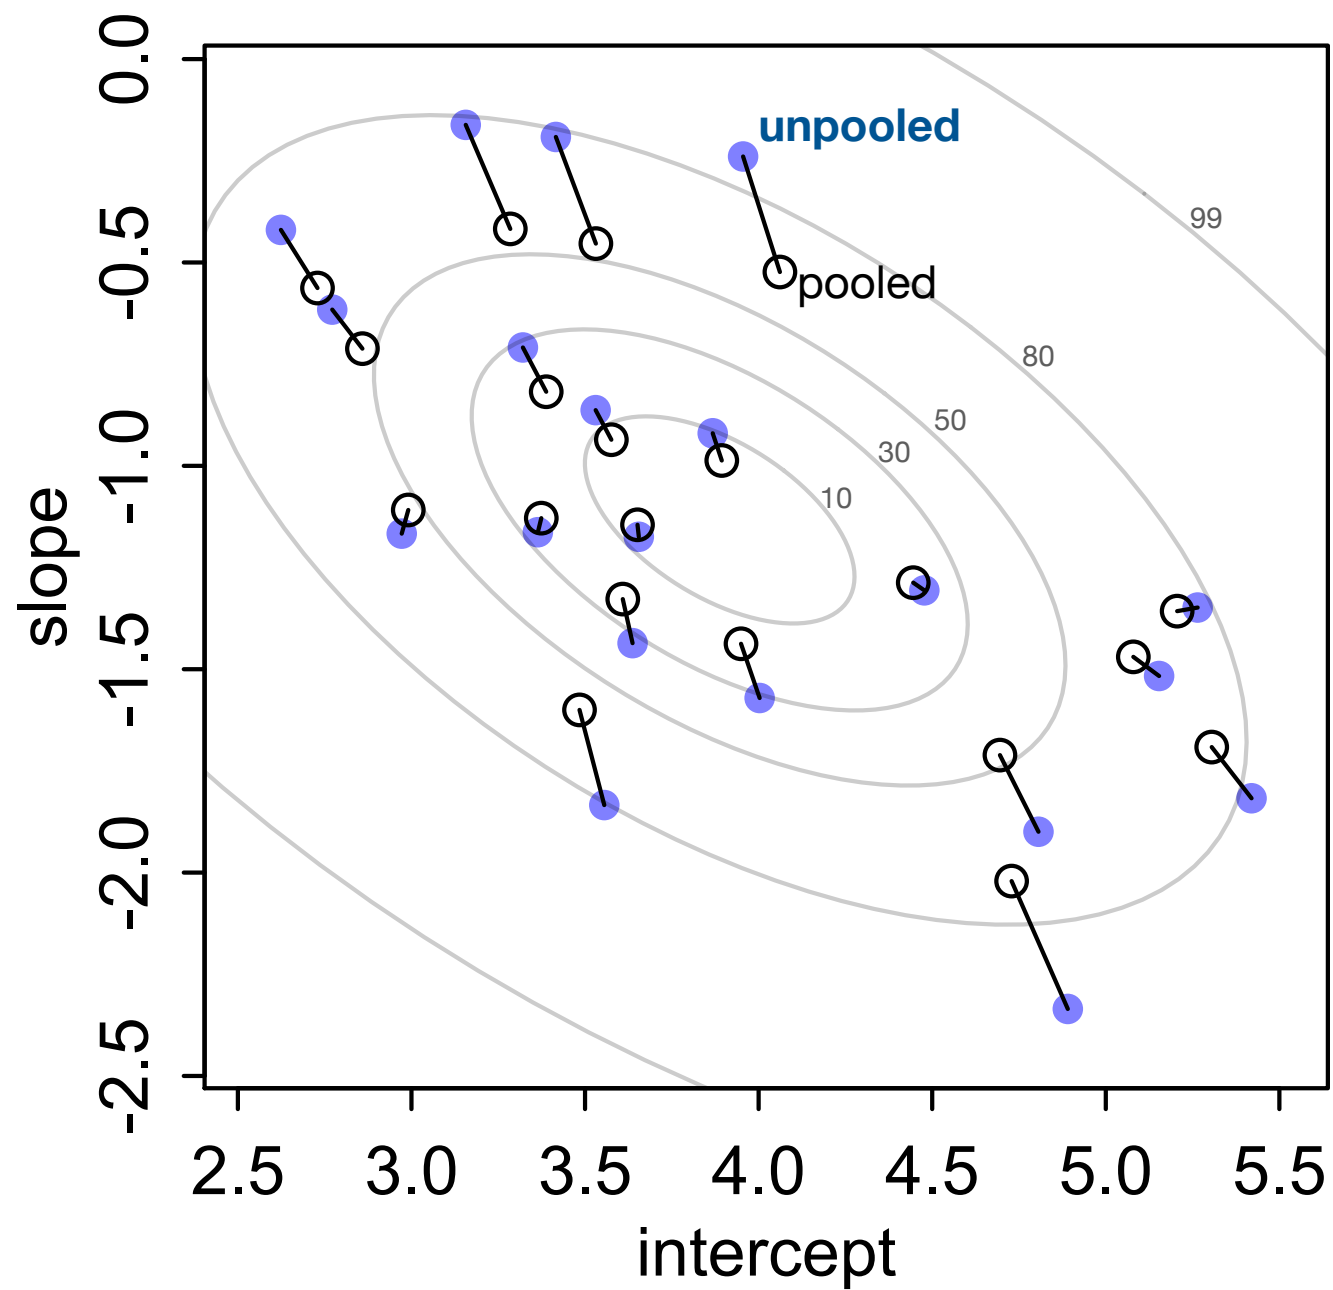
Posterior correlation

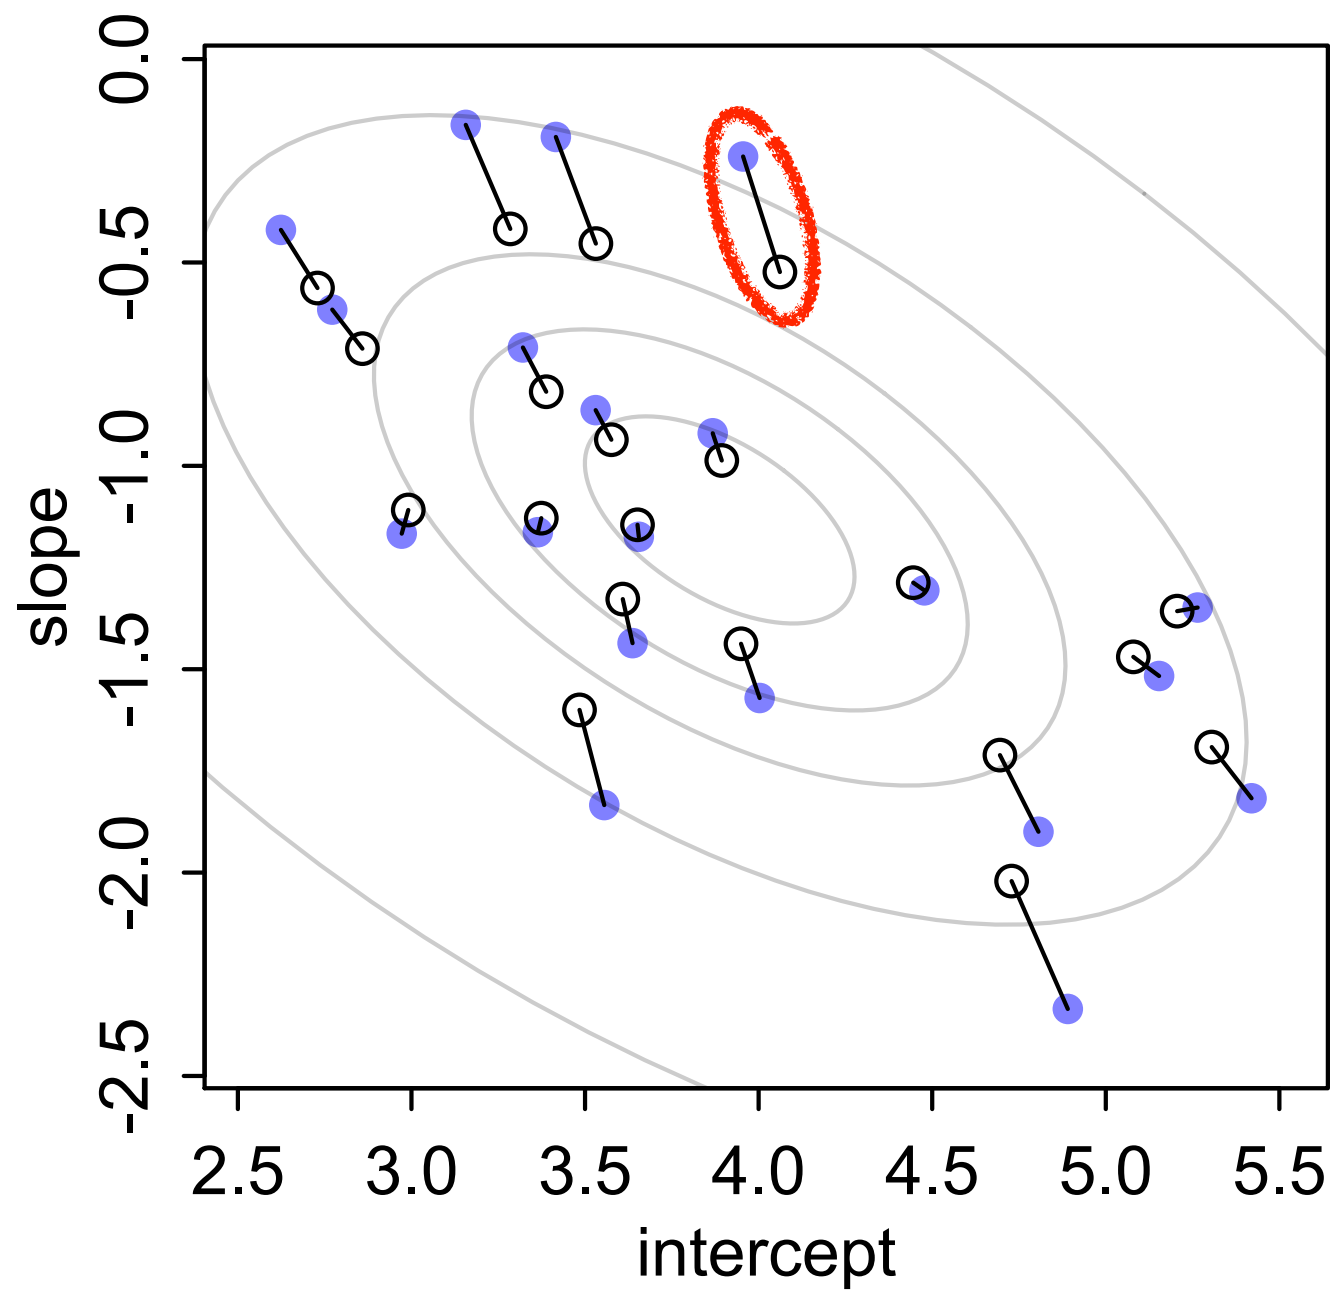
R code
13.12 `post <- extract.samples(m13.1)`
`dens(post$Rho[,1,2])`




Posterior shrinkage



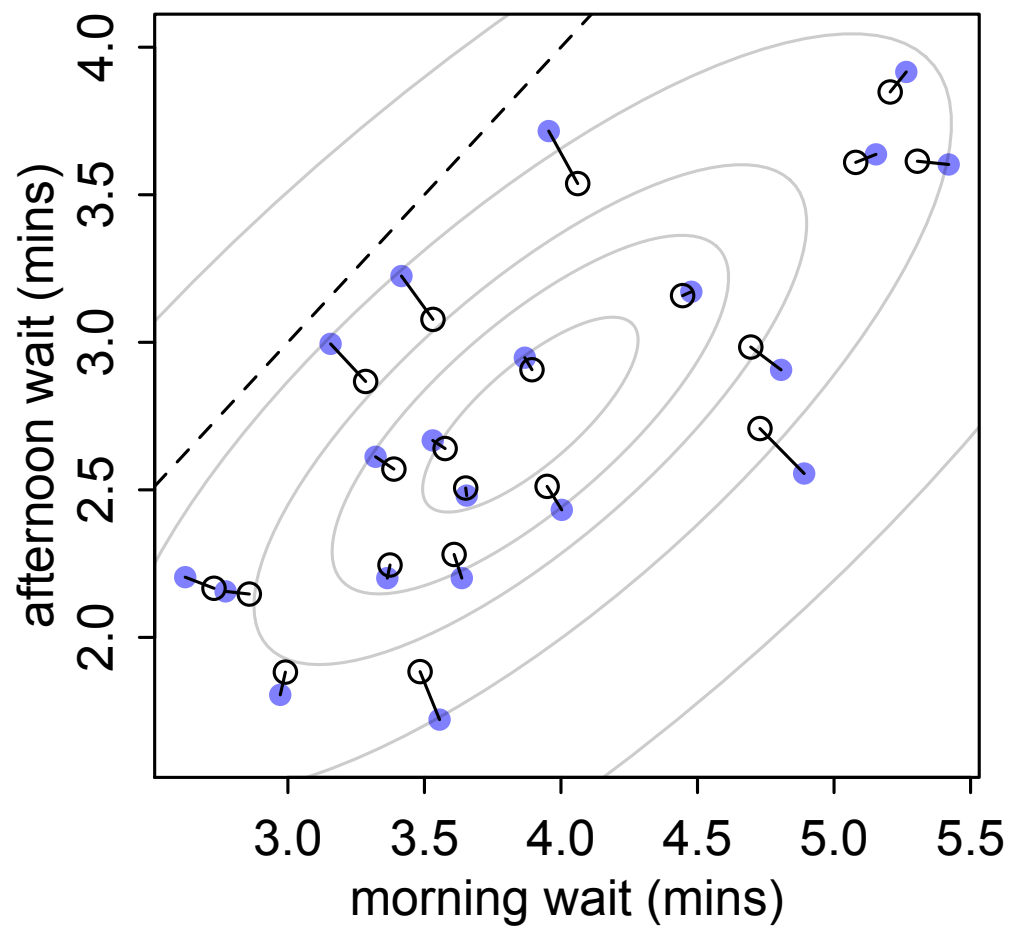




parameter scale




outcome scale



Multi-dimensional shrinkage

- Joint distribution of varying effects pools information across intercepts & slopes
- Correlation btw effects => shrinkage in one dimension induces shrinkage in others
- Improved accuracy, just like varying intercepts



Example: UCB admit data again

	dept	applicant.gender	admit	reject	applications	male	i	j
1	A	male	512	313	825	1	1	1
2	A	female	89	19	108	0	2	1
3	B	male	353	207	560	1	3	2
4	B	female	17	8	25	0	4	2
5	C	male	120	205	325	1	5	3
6	C	female	202	391	593	0	6	3
7	D	male	138	279	417	1	7	4
8	D	female	131	244	375	0	8	4
9	E	male	53	138	191	1	9	5
10	E	female	94	299	393	0	10	5
11	F	male	22	351	373	1	11	6
12	F	female	24	317	341	0	12	6

Varying intercepts by dept

$$A_i \sim \text{Binomial}(n_i, p_i)$$

$$\text{logit}(p_i) = \alpha_{\text{DEPT}[i]} + \beta m_i$$

$$\alpha_{\text{DEPT}} \sim \text{Normal}(\alpha, \sigma)$$

$$\alpha \sim \text{Normal}(0, 10)$$

$$\beta \sim \text{Normal}(0, 1)$$

$$\sigma \sim \text{HalfCauchy}(0, 2)$$

Varying slopes by dept

$$A_i \sim \text{Binomial}(n_i, p_i)$$

$$\text{logit}(p_i) = \alpha_{\text{DEPT}[i]} + \beta_{\text{DEPT}[i]} m_i$$

$$\begin{bmatrix} \alpha_{\text{DEPT}} \\ \beta_{\text{DEPT}} \end{bmatrix} \sim \text{MVNormal} \left(\begin{bmatrix} \alpha \\ \beta \end{bmatrix}, \mathbf{S} \right)$$

$$\mathbf{S} = \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix} \mathbf{R} \begin{pmatrix} \sigma_\alpha & 0 \\ 0 & \sigma_\beta \end{pmatrix}$$

$$\alpha \sim \text{Normal}(0, 10)$$

$$\beta \sim \text{Normal}(0, 1)$$

$$(\sigma_\alpha, \sigma_\beta) \sim \text{HalfCauchy}(0, 2)$$

$$\mathbf{R} \sim \text{LKJcorr}(2)$$